

ps9

March 17, 2019

```
In [1]: import pandas as pd
        drink=pd.read_csv("strongdrink.txt")
        drink.describe()
```

```
Out[1]:
```

	cultivar	alco	malic	ash	alk	magn	\
count	176.000000	176.000000	176.000000	176.000000	176.000000	176.000000	
mean	1.926136	13.006534	2.327159	2.367386	19.492045	99.840909	
std	0.771047	0.814431	1.117747	0.275617	3.355821	14.329499	
min	1.000000	11.030000	0.740000	1.360000	10.600000	70.000000	
25%	1.000000	12.362500	1.597500	2.210000	17.175000	88.000000	
50%	2.000000	13.050000	1.845000	2.360000	19.500000	98.000000	
75%	3.000000	13.682500	3.047500	2.560000	21.500000	107.250000	
max	3.000000	14.830000	5.800000	3.230000	30.000000	162.000000	

	tot_phen	flav	nonfl_phen	proanth	color_int	hue	\
count	176.000000	176.000000	176.000000	176.000000	176.000000	176.000000	
mean	2.298920	2.043352	0.359545	1.597727	5.031761	0.961000	
std	0.627333	0.995579	0.123046	0.571958	2.317965	0.227225	
min	0.980000	0.340000	0.130000	0.410000	1.280000	0.480000	
25%	1.747500	1.242500	0.267500	1.250000	3.200000	0.790000	
50%	2.380000	2.155000	0.340000	1.560000	4.640000	0.975000	
75%	2.800000	2.882500	0.430000	1.952500	6.147500	1.120000	
max	3.880000	5.080000	0.660000	3.580000	13.000000	1.710000	

	OD280rat	proline
count	176.000000	176.000000
mean	2.623409	748.477273
std	0.705369	316.208737
min	1.270000	278.000000
25%	1.990000	500.000000
50%	2.780000	673.500000
75%	3.172500	986.250000
max	4.000000	1680.000000


(a)

```
In [19]: import matplotlib.pyplot as plt
         %matplotlib inline
```

```

for cultivar, group in drink.groupby(['cultivar']):
    plt.scatter(group['alco'], group['color_int'], label="cultivar"+str(cultivar))
plt.legend()
plt.title("Scatterplot of color intensity and alcohol")
plt.show()

```



(b)

```

In [15]: from sklearn.linear_model import LogisticRegression
drink["constant"]=1
y = drink["cultivar"].values
X = drink[["constant","alco","malic","tot_phen",
           "color_int"]].values

from scipy.stats import uniform as sp_uniform
param_dist1 = {"penalty": ["l1", "l2"], "C": sp_uniform(0.1, 10.0)}

from sklearn.model_selection import RandomizedSearchCV
random_search = RandomizedSearchCV(LogisticRegression(), param_distributions=param_dist1,
                                   n_iter=200, n_jobs=-1, cv=5, random_state=25, scoring='f1')
rs_fit1 = random_search.fit(X, y)
print("optimal parameter values =", rs_fit1.best_estimator_)
print("optimal tuning parameter values =", rs_fit1.best_params_)
print("MSE = ", abs(rs_fit1.best_score_))#best_score_ can be negative

optimal parameter values = LogisticRegression(C=2.665871587495725, class_weight=None, dual=False,
fit_intercept=True, intercept_scaling=1, max_iter=100,

```

```

        multi_class='warn', n_jobs=None, penalty='l1', random_state=None,
        solver='warn', tol=0.0001, verbose=0, warm_start=False)
optimal tuning parameter values = {'C': 2.665871587495725, 'penalty': 'l1'}
MSE = 0.11931818181818182

```

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning:
DeprecationWarning)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning:
FutureWarning)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:460: FutureWarning:
"this warning.", FutureWarning)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear
"the number of iterations.", ConvergenceWarning)

```

(c)

```

In [16]: from scipy.stats import randint as sp_randint
        param_dist2 = {"n_estimators": [10, 200],
                        "max_depth": [2, 4],
                        "min_samples_split": sp_randint(2, 20),
                        "min_samples_leaf": sp_randint(2, 20),
                        "max_features": sp_randint(1, 4)}

        from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier

        random_search2 = RandomizedSearchCV(RandomForestRegressor(), param_distributions=param_dist2,
                                              n_iter=200, n_jobs=-1, cv=5, random_state=25, scoring='mse')
        rs_fit2 = random_search2.fit(X, y)
        print("optimal parameter values =", rs_fit2.best_estimator_)
        print("optimal tuning parameter values =", rs_fit2.best_params_)
        print("MSE = ", abs(rs_fit2.best_score_))#best_score_ can be negative

optimal parameter values = RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=4,
max_features=3, max_leaf_nodes=None, min_impurity_decrease=0.0,
min_impurity_split=None, min_samples_leaf=2,
min_samples_split=3, min_weight_fraction_leaf=0.0,
n_estimators=200, n_jobs=None, oob_score=False,
random_state=None, verbose=0, warm_start=False)
optimal tuning parameter values = {'max_depth': 4, 'max_features': 3, 'min_samples_leaf': 2, 'min_samples_split': 3}
MSE = 0.22442108484807236

```

```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning:
DeprecationWarning)

```

(d)

```
In [17]: from sklearn.svm import SVC
param_dist3 = {'C': sp_uniform(loc=0.1, scale=10.0),
               'gamma': ['scale', 'auto'],
               'shrinking': [True, False]}
random_search3 = RandomizedSearchCV(SVC(kernel="rbf"), param_distributions=param_dist3,
                                   n_iter=200, n_jobs=-1, cv=5, random_state=25, scoring='neg_log_loss')
rs_fit3 = random_search3.fit(X, y)
print("optimal parameter values =", rs_fit3.best_estimator_)
print("optimal tuning parameter values =", rs_fit3.best_params_)
print("MSE = ", abs(rs_fit3.best_score_))#best_score_ can be negative

optimal parameter values = SVC(C=3.3605112613782553, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
optimal tuning parameter values = {'C': 3.3605112613782553, 'gamma': 'scale', 'shrinking': True}
MSE = 0.1534090909090909

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning:
    DeprecationWarning)
```

(e)

```
In [18]: from sklearn.neural_network import MLPClassifier
param_dist4 = {"hidden_layer_sizes": sp_randint(1, 100),
               "activation": ["logistic", "relu"],
               "alpha": sp_uniform(0.1, 10.0)}
random_search4 = RandomizedSearchCV(MLPClassifier(), param_distributions=param_dist4,
                                   n_iter=200, n_jobs=-1, cv=5, random_state=25, scoring='neg_log_loss')
rs_fit4 = random_search4.fit(X, y)
print("optimal parameter values =", rs_fit4.best_estimator_)
print("optimal tuning parameter values =", rs_fit4.best_params_)
print("MSE = ", abs(rs_fit4.best_score_))#best_score_ can be negative

optimal parameter values = MLPClassifier(activation='relu', alpha=3.0723443366017835, batch_size='auto',
    beta_1=0.9, beta_2=0.999, early_stopping=False, epsilon=1e-08,
    hidden_layer_sizes=96, learning_rate='constant',
    learning_rate_init=0.001, max_iter=200, momentum=0.9,
    n_iter_no_change=10, nesterovs_momentum=True, power_t=0.5,
    random_state=None, shuffle=True, solver='adam', tol=0.0001,
    validation_fraction=0.1, verbose=False, warm_start=False)
optimal tuning parameter values = {'activation': 'relu', 'alpha': 3.0723443366017835, 'hidden_layer_sizes': 96}
MSE = 0.20454545454545456

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model_selection\_search.py:841: DeprecationWarning:
    DeprecationWarning)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\normal_network\multilayer_perceptron.py:562  
% self.max_iter, ConvergenceWarning)
```

(f)

Comparing the mse of each model, I think the multinomial logistic model is the best predictor of cultivar since its mse is around 0.12, which is the smallest.